
Gaussian Naive Bayes for Multiclass Classification

— Edouard Clocheret, Eric Fan,
Shivam Hingorani and Junhui Huang —

Overview of Gaussian Naive Bayes

Input:

1. Features X :

- **Dimension:** $X \in \mathbb{R}^{n \times d}$, where:
 - n : Number of samples.
 - d : Number of features.

2. Labels y :

- **Dimension:** $y \in \mathbb{R}^n$, where:
 - n : Number of samples.

Output:

Final prediction Class y

Strengths:

1. Simple and Fast
2. Works Well with Small and Large Datasets
3. Works Well with High-Dimensional Data

Mathematical Background

Bayes Theorem

Bayes' theorem provides a way to compute the probability of a class C given some features X :

$$P(C | X) = \frac{P(X | C) \cdot P(C)}{P(X)}$$

- $P(C | X)$: **Posterior probability** — the probability of class C given the features X .
- $P(X | C)$: **Likelihood** — the probability of observing X given that the class is C .
- $P(C)$: **Prior probability** — the probability of class C before observing the data.
- $P(X)$: **Evidence** — the total probability of X across all classes (this is constant for all classes and is often ignored during prediction).

Mathematical background of the algorithm

The goal: Predict the class label based on the features among multiple classes.

$$C = \arg \max_{C_k} P(C_k | X) = \arg \max_{C_k} P(C_k) \cdot \prod_{i=1}^n P(x_i | C_k)$$

First, calculate the prior distribution.

which is defined as the probability of each class based on the training data.

$$P(C_k) = \frac{\text{Number of samples in class } C_k}{\text{Total number of samples}} = \frac{N_k}{N}$$

Second, calculate the likelihood.

which is the probability of observing **feature X** given that the data point belongs to a specific **Class C_k** .

Naive Assumption (features are independent)

$$P(X|C_k) = \prod_{i=1}^n P(x_i|C_k)$$

Detailed look on likelihood (Gaussian distribution)

$$P(x_i|C_k) = \frac{1}{\sqrt{2\pi\sigma_k^2}} \exp\left(-\frac{(x_i - \mu_k)^2}{2\sigma_k^2}\right)$$

Third, calculate the posterior likelihood.

Since $P(X)$ is the same for all classes, we can ignore it when comparing classes.

$$P(C_k | X) \propto P(C_k) \cdot \prod_{i=1}^n P(x_i | C_k)$$

Finally, pick the class with the highest posterior probability

$$C = \arg \max_{C_k} P(C_k | X) = \arg \max_{C_k} P(C_k) \cdot \prod_{i=1}^n P(x_i | C_k)$$

Loss (1/2)

$$\text{Likelihood} = \prod_{i=1}^N P(y_i|X_i) \quad \longrightarrow \quad \text{NLL} = - \sum_{i=1}^N \log P(y_i|X_i)$$

Likelihood across all N samples, that has to be **maximized**.

It represents the **probability** that a label is **correctly assigned**.

Loss = Negative Log Likelihood, that has to be **minimized**.

We use the **log** loss because it is easier to calculate derivatives, and more efficient to **sum log probabilities**.

Loss (2/2)

Assumption: the features follow a Gaussian distribution.

$$P(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

So the log-Loss becomes :

$$\log L(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

Maximum Likelihood Estimators (1/2)

$$P(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

How to compute **estimators** of the parameters μ and σ^2 that **minimizes the Loss?**

Maximum Likelihood Estimators (2/2)

$$\log L(\mu, \sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

Taking the derivative with respect to μ and σ^2 and setting them to zero, we obtain:

- The MLE for the mean:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

The result of MLE is the empirical mean and variance.

- The MLE for the variance:

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

Model

Model Implementation (1/2)

Input:

- Training data X_{train} (N samples, n features)
- Labels y_{train} (corresponding to X_{train})
- Test data X_{test}

Output:

- Predicted labels y_{pred} for X_{test}

Step 1: Compute Class Statistics

- Separate X_{train} by class into subsets $\{X_{C1}, X_{C2}, \dots, X_{Ck}\}$, one for each class.
- For each class C_k :
 - Compute prior probability $P(C_k)$:
$$P(C_k) = N_{Ck} / N_{\text{total}}$$
 - For each feature i :
 - Compute mean $\mu_{Ck,i}$:
$$\mu_{Ck,i} = \text{mean}(X_{Ck}[:, i])$$
 - Compute variance $\sigma_{Ck,i}^2$:
$$\sigma_{Ck,i}^2 = \text{variance}(X_{Ck}[:, i])$$

Model Implementation (2/2)

Step 2: Define Gaussian Probability Density Function

```
Function Gaussian(x,  $\mu$ ,  $\sigma^2$ ):  
    return (1 / sqrt(2 *  $\pi$  *  $\sigma^2$ )) * exp(-((x -  $\mu$ )^2) / (2 *  $\sigma^2$ ))
```

Step 3: Predict Class for Test Data

```
For each sample x in X_test:  
    For each class C_k:  
        - Compute log of posterior probability:  
            log P(C_k | x) = log P(C_k)  
                           +  $\sum$  (log Gaussian(x_i,  $\mu_{C_k,i}$ ,  $\sigma_{C_k,i}^2$ ) for i in features)  
        - Assign x to the class C_k with the highest log P(C_k | x)
```

Step 4: Return Predictions

```
Return y_pred (predicted labels for all samples in X_test)
```

Reproduction

Original Study

- Multiclass Classification with Iris Dataset using Gaussian Naive Bayes, *International Journal of Computer Science and Mobile Computing*
- Uses “the Iris dataset” from UCI Machine Learning repository, included in scikit-learn
- Train a Gaussian Naive Bayes classifier using 80% of data, then test using the remaining 20%

Data

- 150 examples, evenly distributed
- 3 labels / classes: Iris setosa, Iris versicolor, Iris virginica
- 4 features: sepal length, sepal width, petal length, petal width



Iris setosa



Iris versicolor



Iris virginica

Performance

1. Accuracy (A) = $(TP + TN) / (TP + TN + FP + FN)$
2. Precision (P) = $(TP) / (TP + FP)$
3. Recall (R) = $(TP) / (TP + FN)$
4. F1-score = $2 \cdot (P \cdot R) / (P + R)$

Results

Accuracy	Precision	Recall	F1-score
0.95	0.95	0.95	0.95

- Our results closely mirror those produced by previous studies using this dataset and align with the scikit-learn implementation of GNB
- **90% accuracy is often considered 'good', and we are comfortably above that**
- 77% test accuracy on "Student Performance" dataset

Limitations

- Naive assumption (feature independence)
 - Poor Performance with Correlated Features
- Gaussian Distribution Assumption
- Sensitivity to Data Scaling
 - Features with larger variances can dominate the calculation of probabilities

Thank you for your attention!

Edouard Clocheret, Eric Fan, Shivam Hingorani and Junhui Huang

References

1. Wikipedia, 2023. Naive Bayes Classifier. [online] Available at: https://en.wikipedia.org/wiki/Naive_Bayes_classifier [Accessed 8 Dec. 2024].
2. GeeksforGeeks, 2023. Gaussian Naive Bayes. [online] Available at: <https://www.geeksforgeeks.org/gaussian-naive-bayes/> [Accessed 8 Dec. 2024].
3. Brown University, n.d. DATA2060, Assignment 6. [online] Brown University. Available at: [Accessed 8 Dec. 2024].
4. Scikit-learn, 2024. Naive Bayes. [online] Available at: https://scikit-learn.org/1.5/modules/naive_bayes.html [Accessed 8 Dec. 2024].
5. Iqbal, Z. and Yadav, M., 2020. Multiclass Classification with Iris Dataset Using Gaussian Naive Bayes. International Journal of Computer Science and Mobile Computing, 9(4), pp.27–35. Available at: [Link](#) [Accessed 8 Dec. 2024].
6. UCI Machine Learning Repository, n.d. Iris Data Set. [online] Available at: <https://archive.ics.uci.edu/ml/datasets/Iris> [Accessed 8 Dec. 2024].

Appendices

The Model Class

Attributes :

```
self.n_classes  
self.feature_dist  
self.priors
```

```
class GaussNaiveBayes :  
    """ Gaussian Naive Bayes model for multiclass classification  
  
    @attrs:  
        n_classes:    the number of classes  
        feature_dist:  a 3D (n_classes x n_features x 2) NumPy array of  
                        the attribute distributions: mean and sigma2 for  
                        each feature for each class  
  
                        ex : [[mu, sigma2],    of feature 1  
                             [mu, sigma2],    of feature 2  
                             [mu, sigma2]],   of feature 3 of class 1  
  
                             [[mu, sigma2],   of feature 1  
                              [mu, sigma2],   of feature 2  
                              [mu, sigma2]]]   of feature 3 of class 2  
  
        priors: a 1D NumPy array of the priors distribution  
  
                        ex : [p_class1, p_class2, p_class3]  
  
    """
```

Training (1/3)

```
for class_label in range(self.n_classes):  
    X_class = X_train [y_train == class_label]
```

For each class

Training (2/3)

```
for class_label in range(self.n_classes):  
    X_class = X_train [y_train == class_label]  
  
    # Computing prior with Laplace smoothing  
    total_class = X_class.shape[0]  
    a = self.laplace_smoothing  
    self.priors[class_label] = (total_class + a) / (n_examples + a * self.n_classes)
```

Computing of the prior with Laplace Smoothing

Training (3/3)

```
for class_label in range(self.n_classes):  
    X_class = X_train[y_train == class_label]  
  
    # Computing prior with Laplace smoothing  
    total_class = X_class.shape[0]  
    a = self.laplace_smoothing  
    self.priors[class_label] = (total_class + a) / (n_examples + a * self.n_classes)  
  
    # Computing the moments (mean mu and variance sigma2)  
    mu = np.mean(X_class, axis=0)  
    sigma2 = np.var(X_class, axis = 0) + 1e-6
```

Computing and saving the moments

Prediction (1/4)

```
predictions[index] = np.argmax(logprobs)
```

Objective : find the maximum of probability over classes

Prediction (2/4)

```
logprobs = np.log(self.priors)
```

First compute the priors for each class

```
predictions[index] = np.argmax(logprobs)
```


Prediction (3/4)

```
logprobs = np.log(self.priors)

for class_label in range(self.n_classes):
    mu = self.feature_dist[class_label, :, 0] # All mus for each feature at a time
    sigma2 = self.feature_dist[class_label, :, 1]
```

```
predictions[index] = np.argmax(logprobs)
```

Compute the moments for each class

Prediction (4/4)

```
logprobs = np.log(self.priors)

for class_label in range(self.n_classes):
    mu = self.feature_dist[class_label, :, 0] # All mus for each feature at a time
    sigma2 = self.feature_dist[class_label, :, 1]

    logprobs[class_label] += np.sum(np.log(gaussian(input, mu, sigma2) + 1e-6))

predictions[index] = np.argmax(logprobs)
```

The result is the sum of the logprobs
And we choose the argmax

Unit test

```
pri1, feature1 = test_model1.train(x1,y1)
assert (np.array(pri1) == pytest.approx(np.array([0.571,0.428]),0.01))
feature1 = feature1.flatten()
cmp = np.array([[0.00, 0.00], [0.333, 0.222], [0.667, 0.222]], [[1.00, 0.00], [0.50, 0.25], [1.00, 0.00]]).flatten()
assert (feature1 == pytest.approx(cmp, abs = 0.01))

assert test_model1.accuracy(x1,y1) == 1.
assert test_model1.accuracy(x2,y2) >= .8
```

Introduction to dataset

The Iris dataset is a classic dataset in machine learning and statistics, widely used for classification tasks. It comprises 150 samples, equally divided among three species of Iris flowers: Iris setosa, Iris virginica, and Iris versicolor. Each sample includes four numerical features: sepal length, sepal width, petal length, and petal width, measured in centimeters.

```
iris = load_iris()  
X = iris.data  
y = iris.target
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Test result

```
#Train-test on real data
model1 = GaussNaiveBayes()
model1.train(X_train, y_train)
print("-----")
print("Train accuracy:")
m1train = model1.accuracy(X_train, y_train)
print(m1train)#change by GF
print("-----")
print("Test accuracy:")
m1test = model1.accuracy(X_test, y_test)
print(m1test)#change by GF
print("-----")
```

Train accuracy:

0.95

Test accuracy:

1.0
