
Gaussian Naive Bayes for Multiclass Classification

— Edouard Clocheret, Eric Fan,
Shivam Hingorani, Junhui Huang —

Overview of Gaussian Naive Bayes

Input:

1. Features X :

- **Dimension:** $X \in \mathbb{R}^{n \times d}$, where:
 - n : Number of samples.
 - d : Number of features.

2. Labels y :

- **Dimension:** $y \in \mathbb{R}^n$, where:
 - n : Number of samples.

Output:

Final prediction Class y

Strengths:

1. Simple and Fast
2. Works Well with Small and Large Datasets
3. Works Well with High-Dimensional Data

Mathematical Background

Bayes Theorem

Bayes' theorem provides a way to compute the probability of a class C given some features X :

$$P(C | X) = \frac{P(X | C) \cdot P(C)}{P(X)}$$

- $P(C | X)$: **Posterior probability** — the probability of class C given the features X .
- $P(X | C)$: **Likelihood** — the probability of observing X given that the class is C .
- $P(C)$: **Prior probability** — the probability of class C before observing the data.
- $P(X)$: **Evidence** — the total probability of X across all classes (this is constant for all classes and is often ignored during prediction).

Mathematical background of the algorithm

The goal: Predict the class label based on the features among multiple classes.

$$C = \arg \max_{C_k} P(C_k | X) = \arg \max_{C_k} P(C_k) \cdot \prod_{i=1}^n P(x_i | C_k)$$

First, calculate the prior distribution.

which is defined as the probability of each class based on the training data.

$$P(C_k) = \frac{\text{Number of samples in class } C_k}{\text{Total number of samples}} = \frac{N_k}{N}$$

Second, calculate the likelihood.

which is the probability of observing **feature X** given that the data point belongs to a specific **Class C_k** .

Naive Assumption (features are independent)

$$P(X|C_k) = \prod_{i=1}^n P(x_i|C_k)$$

Detailed look on likelihood (Gaussian distribution)

$$P(x_i|C_k) = \frac{1}{\sqrt{2\pi\sigma_k^2}} \exp\left(-\frac{(x_i - \mu_k)^2}{2\sigma_k^2}\right)$$

Third, calculate the posterior likelihood.

Since $P(X)$ is the same for all classes, we can ignore it when comparing classes.

$$P(C_k | X) \propto P(C_k) \cdot \prod_{i=1}^n P(x_i | C_k)$$

Finally, pick the class with the highest posterior probability

$$C = \arg \max_{C_k} P(C_k \mid X) = \arg \max_{C_k} P(C_k) \cdot \prod_{i=1}^n P(x_i \mid C_k)$$

Loss (1/2)

$$\text{Likelihood} = \prod_{i=1}^N P(y_i|X_i) \quad \longrightarrow \quad \text{NLL} = - \sum_{i=1}^N \log P(y_i|X_i)$$

Likelihood across all N samples, that has to be **maximized**.

It represents the **probability** that a label is **correctly assigned**.

Loss = Negative Log Likelihood, that has to be **minimized**.

We use the **log** loss because it is easier to calculate derivatives, and more efficient to **sum** log **probabilities**.

Loss (2/2)

Assumption: the features follow a Gaussian distribution.

$$P(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

So the log-Loss becomes :

$$\log L(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

Maximum Likelihood Estimators (1/2)

$$P(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

How to compute **estimators** of the parameters μ and σ^2 that **minimizes the Loss?**

Maximum Likelihood Estimators (2/2)

$$\log L(\mu, \sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

Taking the derivative with respect to μ and σ^2 and setting them to zero, we obtain:

- The MLE for the mean:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

The result of MLE is the empirical mean and variance.

- The MLE for the variance:

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

Model

Model Implementation (1/2)

Input:

- Training data X_{train} (N samples, n features)
- Labels y_{train} (corresponding to X_{train})
- Test data X_{test}

Output:

- Predicted labels y_{pred} for X_{test}

Step 1: Compute Class Statistics

- Separate X_{train} by class into subsets $\{X_{C1}, X_{C2}, \dots, X_{Ck}\}$, one for each class.
- For each class C_k :
 - Compute prior probability $P(C_k)$:
$$P(C_k) = N_{Ck} / N_{\text{total}}$$
 - For each feature i :
 - Compute mean $\mu_{Ck,i}$:
$$\mu_{Ck,i} = \text{mean}(X_{Ck}[:, i])$$
 - Compute variance $\sigma_{Ck,i}^2$:
$$\sigma_{Ck,i}^2 = \text{variance}(X_{Ck}[:, i])$$

Model Implementation (2/2)

Step 2: Define Gaussian Probability Density Function

```
Function Gaussian(x,  $\mu$ ,  $\sigma^2$ ):  
    return (1 / sqrt(2 *  $\pi$  *  $\sigma^2$ )) * exp(-((x -  $\mu$ )^2) / (2 *  $\sigma^2$ ))
```

Step 3: Predict Class for Test Data

```
For each sample x in X_test:  
    For each class C_k:  
        - Compute log of posterior probability:  
            log P(C_k | x) = log P(C_k)  
                        +  $\sum$  (log Gaussian(x_i,  $\mu_{C_k,i}$ ,  $\sigma_{C_k,i}^2$ ) for i in features)  
        - Assign x to the class C_k with the highest log P(C_k | x)
```

Step 4: Return Predictions

```
Return y_pred (predicted labels for all samples in X_test)
```

Reproduction

Original Study

- Multiclass Classification with Iris Dataset using Gaussian Naive Bayes, *International Journal of Computer Science and Mobile Computing*
- Uses “the Iris dataset” from UCI Machine Learning repository, included in scikit-learn
- Train a Gaussian Naive Bayes classifier using 80% of data, then test using the remaining 20%

Data

- 150 examples, evenly distributed
- 3 labels / classes: Iris setosa, Iris versicolor, Iris virginica
- 4 features: sepal length, sepal width, petal length, petal width



Iris setosa



Iris versicolor



Iris virginica

Performance

1. Accuracy (A) = $(TP + TN) / (TP + TN + FP + FN)$
2. Precision (P) = $(TP) / (TP + FP)$
3. Recall (R) = $(TP) / (TP + FN)$
4. F1-score = $2 \cdot (P \cdot R) / (P + R)$

Results

Accuracy	Precision	Recall	F1-score
0.95	0.95	0.95	0.95

- Our results closely mirror those produced by previous studies using this dataset and align with the scikit-learn implementation of GNB
- **90% accuracy is often considered 'good', and we are comfortably above that**
- 77% test accuracy on "Student Performance" dataset

Limitations

- Naive assumption (feature independence)
 - Poor Performance with Correlated Features
- Gaussian Distribution Assumption
- Sensitivity to Data Scaling
 - Features with larger variances can dominate the calculation of probabilities

Thank you for your attention!

Edouard Clocheret, Eric Fan, Shivam Hingorani, Junhui Huang

References

1. Wikipedia, 2023. Naive Bayes Classifier. [online] Available at: https://en.wikipedia.org/wiki/Naive_Bayes_classifier [Accessed 8 Dec. 2024].
2. GeeksforGeeks, 2023. Gaussian Naive Bayes. [online] Available at: <https://www.geeksforgeeks.org/gaussian-naive-bayes/> [Accessed 8 Dec. 2024].
3. Brown University, n.d. DATA2060, Assignment 6. [online] Brown University. Available at: [Accessed 8 Dec. 2024].
4. Scikit-learn, 2024. Naive Bayes. [online] Available at: https://scikit-learn.org/1.5/modules/naive_bayes.html [Accessed 8 Dec. 2024].
5. Iqbal, Z. and Yadav, M., 2020. Multiclass Classification with Iris Dataset Using Gaussian Naive Bayes. International Journal of Computer Science and Mobile Computing, 9(4), pp.27–35. Available at: [Link](#) [Accessed 8 Dec. 2024].
6. UCI Machine Learning Repository, n.d. Iris Data Set. [online] Available at: <https://archive.ics.uci.edu/ml/datasets/Iris> [Accessed 8 Dec. 2024].