

final_project

December 13, 2024

1 DATA 2060 – *Final Project*

1.0.1 Machine Learning - from theory to algorithms

Edouard Clocheret, Junhui Huang, Eric Fan, Shivam Hingorani

Link to the github repository: <https://github.com/edouardclocheret/GaussianNaiveBayes.git>

1.1 1. Overview of Gaussian Naive Bayes for classification

Gaussian Naive Bayes is an extension of the conventional Naive Bayes algorithm. Like the original, the Gaussian variety of Naive Bayes is used for classification tasks, specifically those where the features are continuous and can be assumed to follow a Gaussian (i.e., normal) distribution. This means Gaussian Naive Bayes maintains the assumption of conditional independence across features (hence the description as “naive”). This algorithm is computationally efficient, which can be advantageous when using large datasets, and it serves as the basis for many more complex machine learning models/techniques.

1.1.1 REPRESENTATION

Multiclass Naive Bayes Representation: For a given input: $X = (x_1, x_2, \dots, x_n)$,

First, we calculate the prior probability for each class, which is defined as the probability of each class based on the training data. This is denoted as

$$P(C_k) = \frac{\text{Number of samples in class } C_k}{\text{Total number of samples}} = \frac{N_k}{N}$$

Second, For each class, we calculate the likelihood $P(X | C_k)$, which is the probability of observing feature X given that the data point belongs to a specific Class C_k

To calculate this, we use the Naive Bayes assumption that each feature x_i is independent of others given the class C . This allows us to calculate the product of the individual feature likelihoods:

$$P(X|C_k) = \prod_{i=1}^n P(x_i|C_k)$$

Third, We calculate the posterior Probability $P(C_k|X)$ for each class, which is the goal of this algorithm, calculating the probability of each class C based on the feature X . Since $P(X)$

is the same for all classes, we can ignore it when comparing classes, so we don't need to calculate it explicitly.

$$P(C_k|X) \propto P(C_k) \cdot \prod_{i=1}^n P(x_i|C_k)$$

where

$$P(x_i|C_k) = \frac{1}{\sqrt{2\pi\sigma_k^2}} \exp\left(-\frac{(x_i - \mu_k)^2}{2\sigma_k^2}\right)$$

Last, We select the class C with the highest posterior probability, which means that we After computing $P(C_k | X)$ for each class C_k , choose the class with the highest posterior probability as the predicted class:

$$C = \arg \max_{C_k} P(C_k | X) = \arg \max_{C_k} P(C_k) \cdot \prod_{i=1}^n P(x_i|C_k)$$

1.1.2 LOSS

For Gaussian Naive Bayes (GNB), the loss function normally not defined or refined. Because this type of model deduct based on **Maximum Likelihood Estimation** (MLE).

However, if we have to define a loss function when during the Training or Evaluation Process, we can define *Negative Log Likelihood* (NLL) to measure the classification loss.

The formula:

$$\text{NLL} = - \sum_{i=1}^N \log P(y_i|X_i)$$

y_i is the i th sample's true label, X_i is input sample feature, $P(y_i|X_i)$ is the probability of model predict.

Since Gaussian Naive Bayes assumes that features follow a Gaussian (normal) distribution, it calculates the conditional probability using the probability density function of a Gaussian distribution. For each class y and feature x_i , the conditional probability ($P(x_i | y)$) is given by:

$$P(x_i|y) = \frac{1}{\sqrt{2\pi\sigma_y^2}} \exp\left(-\frac{(x_i - \mu_y)^2}{2\sigma_y^2}\right)$$

Such that the negative log likelihood for Gaussian Naive Bayes is:

$$\text{NLL}_{\text{Gaussian}} = - \sum_{i=1}^N \left(\log P(y_i) + \sum_{j=1}^d \log \frac{1}{\sqrt{2\pi\sigma_{y_i,j}^2}} - \frac{(X_{ij} - \mu_{y_i,j})^2}{2\sigma_{y_i,j}^2} \right)$$

$\{y_{i,j}\}$ and $\{\mu_{i,j}\}$ are sample j 's mean and variance under label y_i

Maximum Likelihood Estimation Suppose we have a set of data points $X = \{x_1, x_2, \dots, x_n\}$, and these data points are drawn from a normal distribution with mean μ and variance σ^2 . We want to know μ and σ^2 by using MLE.

The PDF of the normal distribution is:

$$P(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

For a given dataset $X = \{x_1, x_2, \dots, x_n\}$, the joint likelihood function is:

$$L(\mu, \sigma^2) = \prod_{i=1}^n P(x_i|\mu, \sigma^2)$$

Taking the log of the likelihood function, we get the log-likelihood function:

$$\log L(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

Taking the derivative with respect to μ and σ^2 and setting them to zero, we obtain:

- The MLE for the mean:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

- The MLE for the variance:

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

1.1.3 OPTIMIZER

The Gaussian assumption of the model is to consider the probabilities $P(x_i|C_k)$ to follow a Gaussian distribution:

$$P(x_i|C_k) = \frac{1}{\sqrt{2\pi\sigma_C^2}} \exp\left(-\frac{(x_i - \mu_C)^2}{2\sigma_C^2}\right)$$

The mean μ and variance σ estimators are derived from maximum likelihood estimation (MLE).

We use the log-loss because it is easier to calculate derivatives of a sum:

$$\text{LogLoss}(C, X) = -\log P(C) - \sum_{i=1}^n \log\left(\frac{1}{\sqrt{2\pi\sigma_{C,i}^2}} \exp\left(-\frac{(x_i - \mu_{C,i})^2}{2\sigma_{C,i}^2}\right)\right)$$

and simplifies as :

$$\text{LogLoss}(C, X) = -\log P(C) + \sum_{i=1}^n \left(\frac{(x_i - \mu_{C,i})^2}{2\sigma_{C,i}^2} + \log \sqrt{2\pi\sigma_{C,i}^2} \right)$$

We maximize the likelihood (= minimize the loss) by taking the logarithm (log-likelihood) and setting the derivatives with respect to $\mu_{C,i}$ and $\sigma_{C,i}$ to zero, we obtain:

$$\mu_{C_k,i} = \frac{1}{N_C} \sum_{j=1}^{N_C} x_i^{(j)}$$

and

$$\sigma_{C_k,i}^2 = \frac{1}{N_C} \sum_{j=1}^{N_C} (x_i^{(j)} - \mu_{C_k,i})^2$$

where * N_C is the number of samples in class C. * $x_i^{(j)}$ is the value of feature x_i for the j -th sample in class C. —

In practice, we first separate the data by class and use the optimal estimators above for the mean and variance.

The training consists of computing and memorizing $\mu_{C_k,i}$ and $\sigma_{C_k,i}^2$ for all feature x_i and classes C_k .

These values are then used for the classifications to compute $P(x_i|C_k)$.

1.2 2. Model

```
[1]: import numpy as np

def gaussian(x, mu, sigma2):

    sigma2 = np.where(sigma2 == 0, 1e-6, sigma2)    # Change by GF
    coeff = 1 / np.sqrt(2 * np.pi * sigma2)
    exponent = np.exp(-((x - mu) ** 2) / (2 * sigma2))

    return coeff * exponent

class GaussNaiveBayes:
    """ Gaussian Naive Bayes model for multiclass classification

    @attrs:
        n_classes:    the number of classes
        feature_dist:  a 3D (n_classes x n_features x 2) NumPy array of
                        the attribute distributions: mean and sigma2 for
                        each feature for each class
```

```

        ex : [[mu, sigma2],    of feature 1
              [mu, sigma2],    of feature 2
              [mu, sigma2]],   of feature 3 of class 1

              [[mu, sigma2],    of feature 1
              [mu, sigma2],    of feature 2
              [mu, sigma2]]]   of feature 3 of class 2

    priors: a 1D NumPy array of the priors distribution

        ex : [p_class1, p_class2, p_class3]
    """

def __init__(self):
    self.n_classes = None    # Computed at training
    self.feature_dist = None
    self.priors = None

    self.laplace_smoothing = 1 # Default value for the smoothing parameter

def train(self, X_train, y_train):
    """ Trains the model, using maximum likelihood estimation.
    @params:
        X_train: a 2D (n_examples x n_features) numpy array
        y_train: a 1D (n_examples) numpy array of the corresponding labels
    @return:
        self.priors
        self.feature_dist
    """ # Change by GF

    self.n_classes = len(set(y_train))
    n_features = X_train.shape[1]

    self.priors = np.zeros(self.n_classes)
    self.feature_dist = np.zeros((self.n_classes, n_features, 2))

    n_examples = X_train.shape[0]

    for class_label in range(self.n_classes):
        X_class = X_train [y_train == class_label]

        # Computing prior with Laplace smoothing
        total_class = X_class.shape[0]
        a = self.laplace_smoothing
        self.priors[class_label] = (total_class + a) / (n_examples + a * self.
↪n_classes)

```

```

        # Computing the moments (mean mu and variance sigma2)
        mu = np.mean(X_class, axis=0)
        sigma2 = np.var(X_class, axis = 0) + 1e-6

        # Saving the mu and sigma for later predictions
        self.feature_dist[class_label, :, :] = np.vstack((mu, sigma2)).T
    # Change by GF

    return self.priors, self.feature_dist

def predict(self, inputs):
    """ Outputs a predicted label for each input in inputs.

    @params:
        inputs: a 2D NumPy array containing inputs (n_requests * n_features)
    @return:
        a 1D numpy array of predictions
    """

    predictions = np.zeros(len(inputs))

    for index, input in enumerate(inputs) :

        logprobs = np.log(self.priors)

        for class_label in range(self.n_classes):
            mu = self.feature_dist[class_label, :, 0] # All mus for each
            # feature at a time
            sigma2 = self.feature_dist[class_label, :, 1]

            logprobs[class_label] += np.sum(np.log(gaussian(input, mu,
            # sigma2) + 1e-6))

        predictions[index] = np.argmax(logprobs)

    return predictions

def accuracy(self, X_test, y_test):
    """ Outputs the accuracy of the trained model on a given dataset (data).

    @params:
        X_test: a 2D numpy array of examples
        y_test: a 1D numpy array of labels
    @return:

```

```

        a float number indicating accuracy (between 0 and 1)
    """

```

```

    y_pred = self.predict(X_test)
    return np.mean(y_pred == y_test)

```

1.3 3. Check Model

```

[2]: import pytest

test_model1 = GaussNaiveBayes()

x1 = np.array([[0,0,1], [0,1,0], [1,0,1], [1,1,1], [0,0,1]])
y1 = np.array([0,0,1,1,0])
x_test1 = np.array([[1,0,0],[0,0,0],[1,1,1],[0,1,0], [1,1,0]])
y_test1 = np.array([0,0,1,0,1])

x2 = np.array([[0,0,1], [0,1,1], [1,1,1], [1,1,1], [0,0,0], [1,1,0]])
y2 = np.array([0,1,1,1,0,1])
x_test2 = np.array([[0,0,1], [0,1,1], [1,1,1], [1,0,0]])
y_test2 = np.array([0,1,1,0])

pri1, feature1 = test_model1.train(x1,y1)
assert (np.array(pri1) == pytest.approx(np.array([0.571,0.428]),0.01)) # Check
    ↳if the train function return prior have correct value and shape
feature1 = feature1.flatten()
cmp = np.array([[0.00, 0.00], [0.333, 0.222], [0.667, 0.222]], [[1.00, 0.00],
    ↳[0.50, 0.25], [1.00, 0.00]]).flatten()
assert (feature1 == pytest.approx(cmp, abs = 0.01)) # Check if the train
    ↳function return feature have correct value and shape

assert test_model1.accuracy(x1,y1) == 1.    # Check the accuracy function
assert test_model1.accuracy(x2,y2) >= .8
print("test complete")

```

test complete

1.3.1 Recreating previous work

```

[3]: #Load and preprocess data
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

'''

```

The data set consists of 50 samples from each of three species of Iris (Iris_↵
 ↵setosa, Iris virginica and Iris versicolor).
 Four features were measured from each sample: the length and the width of the_↵
 ↵sepals and petals, in centimeters.
 '''
 iris = load_iris()
 X = iris.data
 y = iris.target

 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,↵
 ↵random_state=42)

```
[4]: # Train-test on real data
model1 = GaussNaiveBayes()
model1.train(X_train, y_train)

print("-----")

print("Train accuracy:")
mltrain = model1.accuracy(X_train, y_train)
print(mltrain)  # Change by GF

print("-----")

print("Test accuracy:")
mltest = model1.accuracy(X_test, y_test)
print(mltest)  # Change by GF

print("-----")
```

```
-----
Train accuracy:
0.95
-----
```

```
Test accuracy:
1.0
-----
```

```
[5]: # Comparison with scikit-learn
'''
Use the sklearn model to test whether our model is correct or not.
'''
model = GaussianNB()
model.fit(X_train, y_train)

train_accuracy = accuracy_score(y_train, model.predict(X_train))
test_accuracy = accuracy_score(y_test, model.predict(X_test))
```



```
# If conditional statement executes, our model had the exact same accuracy as
↳ the sklearn model
if train_accuracy == mtrain and test_accuracy == mtest:
    print("Mission Complete")
```

Mission Complete

Comparing results In the previous work we found, which was a journal paper from International Journal of Computer Science and Mobile Computing, they apply the Gaussian Naive Bayes algorithm to classify the Iris dataset into its three species: Iris setosa, Iris versicolor, and Iris virginica.

The dataset consists of 150 samples, with 50 samples per species, and includes four features for classification: sepal length (cm), sepal width (cm), petal length (cm), and petal width (cm). The goal is to utilize these features to predict the class of a flower. They also used scatter plot matrix as part of their EDA to study the relationships between pairs of features.

In their results, they evaluated the performance of the Gaussian Naive Bayes classifier using metrics such as precision (0.95), recall (0.95), F1-score (0.95), and accuracy (0.95). Additionally, they used confusion matrix to explain the performance of the Gaussian Naive Bayes classifier.

The results from our algorithm closely match their findings and are also consistent with the outcomes produced by the built-in scikit-learn implementation.

1.4 4: Github repo

Link to the github repository: <https://github.com/edouardclocheret/GaussianNaiveBayes.git>

1.5 5. References

1. Wikipedia, 2023. Naive Bayes Classifier. [online] Available at: https://en.wikipedia.org/wiki/Naive_Bayes_classifier [Accessed 8 Dec. 2024].
2. GeeksforGeeks, 2023. Gaussian Naive Bayes. [online] Available at: <https://www.geeksforgeeks.org/gaussian-naive-bayes/> [Accessed 8 Dec. 2024].
3. Brown University, n.d. DATA2060, Assignment 6. [online] Brown University. Available at: [Accessed 8 Dec. 2024].
4. Scikit-learn, 2024. Naive Bayes. [online] Available at: https://scikit-learn.org/1.5/modules/naive_bayes.html [Accessed 8 Dec. 2024].
5. Iqbal, Z. and Yadav, M., 2020. Multiclass Classification with Iris Dataset Using Gaussian Naive Bayes. International Journal of Computer Science and Mobile Computing, 9(4), pp.27–35. Available at: [Link](#) [Accessed 8 Dec. 2024].
6. UCI Machine Learning Repository, n.d. Iris Data Set. [online] Available at: <https://archive.ics.uci.edu/ml/datasets/Iris> [Accessed 8 Dec. 2024].

1.6 6. Contributions

Edouard: Coded the core of the model (`train()`, `predict()`, `accuracy()`) and worked on the theory and likelihood maximization.

Eric: Developed unit tests and worked on the theory (loss and optimizer).

Huang: Found previous work where Naive Bayes is applied on Iris dataset. Discussion of previous work, and reproduction of the previous work. Introduced the representation of the Gaussian Naive Bayes.

Shivam: Developed replication of previous work and analyzed results. Researched limitations of Gaussian Naive Bayes Model.

1.7 APPENDIX: Student Performance Dataset

```
[ ]: from google.colab import drive

import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder

drive.mount('/content/drive')
```

```
[ ]: #Import kaggle dataset and show training and testing error.
file_path = '/content/drive/My Drive/Student_performance_data.csv'
df = pd.read_csv(file_path)

y = df['GradeClass']
X = df.drop(columns=['GradeClass', 'StudentID'])

minmax_ftrs = ['Age', 'StudyTimeWeekly', 'Absences']
std_ftrs = ['GPA']

preprocessor = ColumnTransformer(
    transformers=[
        ('minmax', MinMaxScaler(), minmax_ftrs),
        ('std', StandardScaler(), std_ftrs)
    ]
)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
    random_state=42)
pipeline = Pipeline([
```

```

        ('preprocessor', preprocessor),
        ('model', GaussianNB())
    ])

    pipeline.fit(X_train, y_train)

    train_accuracy = pipeline.score(X_train, y_train)
    test_accuracy = pipeline.score(X_test, y_test)

    print("-----")
    print(f"Train Accuracy: {train_accuracy}")

    print("-----")
    print(f"Test Accuracy: {test_accuracy}")

    print("-----")

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
-----
Train Accuracy: 0.7955390334572491
-----
```

```
Test Accuracy: 0.8
-----
```

```

[ ]: file_path = '/content/drive/My Drive/Student_performance_data.csv'
df = pd.read_csv(file_path)
y = df['GradeClass']
X = df.drop(columns=['GradeClass', 'StudentID'])

# Standardize numerical features
numeric_features = ['Age', 'StudyTimeWeekly', 'Absences', 'GPA']
scaler = StandardScaler()
X[numeric_features] = scaler.fit_transform(X[numeric_features])

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(
    X, y_encoded, test_size=0.1, random_state=42, stratify=y_encoded
)

X_train_np = X_train.values
X_test_np = X_test.values
y_train_np = np.array(y_train)
y_test_np = np.array(y_test)

```

```

# Initialize and train the GaussNaiveBayes model
model1 = GaussNaiveBayes()
model1.train(X_train_np, y_train_np)

# Calculate train and test accuracies
print("-----")
mltrain = model1.accuracy(X_train_np, y_train_np)
print(f"Test Accuracy: {mltrain}")
print("-----")
print(f"Test Accuracy: {mltest}")
print("-----")

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
-----
Test Accuracy: 0.7978624535315985
-----
```

```
-----
Test Accuracy: 0.7708333333333334
-----
```